

Name: Muhammad Ihtesham Yousaf

RegN0: CS120202069

Sem: 6th Section B

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import math
```

```
In [4]: data = pd.read_csv('diabetes.csv')
```

```
In [6]: data.head(10)
```

```
Out[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
5	5	116	74	0	0	25.6	0.201
6	3	78	50	32	88	31.0	0.248
7	10	115	0	0	0	35.3	0.134
8	2	197	70	45	543	30.5	0.158
9	8	125	96	0	0	0.0	0.232

```
In [7]: data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype  
---  -
0   Pregnancies                          768 non-null   int64  
1   Glucose                             768 non-null   int64  
2   BloodPressure                       768 non-null   int64  
3   SkinThickness                       768 non-null   int64  
4   Insulin                             768 non-null   int64  
5   BMI                                 768 non-null   float64 
6   DiabetesPedigreeFunction            768 non-null   float64 
7   Age                                 768 non-null   int64  
8   Outcome                             768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [8]: print(data.index)
print((len(data.index)))
```

```
RangeIndex(start=0, stop=768, step=1)
768
```


```
In [9]: sns.countplot(x='Outcome', data=data)
```

```
Out[9]: <AxesSubplot:xlabel='Outcome', ylabel='count'>
```



```
In [11]: data['Age'].plot.hist()
```

```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



```
In [12]: data.isnull()
```

```
Out[12]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunct
0	False	False	False	False	False	False	Fi
1	False	False	False	False	False	False	Fi
2	False	False	False	False	False	False	Fi
3	False	False	False	False	False	False	Fi
4	False	False	False	False	False	False	Fi
...
763	False	False	False	False	False	False	Fi
764	False	False	False	False	False	False	Fi
765	False	False	False	False	False	False	Fi
766	False	False	False	False	False	False	Fi
767	False	False	False	False	False	False	Fi


768 rows × 9 columns

```
In [13]: data.isnull().sum()
```

```
Out[13]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [14]: sns.heatmap(data.isnull())
```

```
Out[14]: <AxesSubplot:>
```



```
In [15]: data.dropna(inplace=True)

In [19]: from sklearn.model_selection import train_test_split

In [21]: X = data.drop(['Outcome'], axis = 1)
y = data['Outcome']

In [22]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, rand

In [23]: from sklearn.linear_model import LogisticRegression

In [24]: classifier = LogisticRegression()

classifier.fit(X_train, y_train)

In [26]: prediction = classifier.predict(X_test)

In [27]: print(prediction)

[0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 1 0 0 0 0 1 0 1 0
0 0 1 0 0 0 0 0 0 0 1 1 0 0 0 0 1 0 1 0 1 0 1 0 1 0 0 0 0 0 0 1 1 1 1 1 1 0
1 0 1 0 0 1 1 0 0 0 0 1 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1 1 0 0 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 1 0 0 1 0 0 1 1 0 0 0
0 0 0 1 0 0]
```

```
In [28]: from sklearn.metrics import accuracy_score

In [29]: accuracy_score(y_test, prediction)

Out[29]: 0.7792207792207793

In [30]: from sklearn.metrics import classification_report

In [33]: print(classification_report(y_test, prediction))

              precision    recall  f1-score   support

     0       0.79        0.90        0.84         99
     1       0.76        0.56        0.65         55

 accuracy          0.78         0.78         0.78        154
 macro avg         0.77         0.73         0.74        154
 weighted avg         0.78         0.78         0.77        154
```

```
In [34]: from sklearn.metrics import confusion_matrix

In [35]: print(confusion_matrix(y_test, prediction))

[[89 10]
 [24 31]]
```